

CE/CZ4045 Natural Language Processing

Regular Expression and Automata (Chapter 2)



Outline

- Regular expression
- How to implement regular expression?
- Finite-state-automata (FSA)
 - Deterministic FSA
 - Non-deterministic FSA



Regular expression

- Regular expression is a **compact textual representation** of a set of strings representing a language
 - In the simplest case, regular expressions describe regular languages
- Have you tried text search with
 - Programming language: perl, python, C#, java....
 - Linux less, emacs, vi, grep, etc..



REs in programming languages

- Java – Pattern

- // Extract the text between the two title elements
- `pattern = "(?i)(<title.*?>)(.+?)(</title>)"`;
- `String updated = EXAMPLE_TEST.replaceAll(pattern, "$2");`
- Reference:
<https://docs.oracle.com/javase/8/docs/api/java/util/regex/Pattern.html>

- Python – re

- `>>> import re`
- `>>> m = re.search('(?!<=abc)def', 'abcdef')`
- `>>> m.group(0)`
- `'def'`



Basic regular expression

RE	Example Patterns Matched
/woodchucks/	“interesting links to <u>woodchucks</u> and lemurs”
/a/	“Ma <u>r</u> y Ann stopped by Mona’s”
/Claire_says,/	“ “Dagmar, my gift please,” <u>Claire</u> says,”
/DOROTHY/	“SURRENDER <u>DOROTHY</u> ”
/!/	“You’ve left the burglar behind again <u>!</u> ” said Nori

- The simplest RE is a sequence of characters
- // are NOT part of the RE, but a notation used in Perl language
- The first instance of each match to the RE is underlined
 - an application might return more than just the first one.



A bit more regular expression

RE	Match	Example Patterns
/[wW]oodchuck/	Woodchuck or woodchuck	“ <u>Wood</u> chuck”
/[abc]/	‘a’, ‘b’, <i>or</i> ‘c’	“In uomini, in sold <u>a</u> ti”
/[1234567890]/	any digit	“plenty of <u>7</u> to 5”

RE	Match	Example Patterns Matched
/[A-Z]/	an upper case letter	“we should call it ‘ <u>D</u> renched Blossoms’ ”
/[a-z]/	a lower case letter	“ <u>m</u> y beans were impatient to be hoed!”
/[0-9]/	a single digit	“Chapter <u>1</u> : Down the Rabbit Hole”

- Square brackets([]) matches **any single character** from within the class
 - [a-zA-Z] matches any letter,
 - [0-9] matches any number;
 - dash (-) specifies a range
- \b matches a **word boundary**
 - /\bdog\b/ The doggie plays in the yard. The dog is cute



A bit more (read Chapter 2)

- Dot (.): matches **any single character** (but “\.” will match a single dot)
 - /3.14159/ vs. /3\.14159/ matches 3.14159 only
- Star (*): matches **zero or more** of preceding item
 - /fred*/ matches fre, fred, fredddd.
- Plus(+): matches **one or more** of preceding item
 - /fred+/ matches fred, fredddd but not fre
- Question mark (?) matches **zero or one** of preceding item

Always match the largest string they can

e.g. /fred*/

String: fredddfff

RE	Match	Example Patterns Matched
woodchucks?	woodchuck or woodchucks	“ <u>woodchuck</u> ”
colou?r	color or colour	“ <u>colour</u> ”

Applications of RE

- Recognize phone numbers
 - e.g. 6790 6300
- Syntax highlighting
 - e.g. html tags, programming language syntax
- Data validation
 - e.g. input for postal code
- Morphological analysis (next chapter)
- And many more ...



Example: When RE is used in real applications

- Find all the instances of the word “the” in a text.
 - `/the/` How about: The?
 - `/[tT]he/` How about: They they other?
 - `/\b[tT]he\b/` How about: the_ the25
- The process we just went through was to fix **two kinds of errors**
 - Matching strings that we should not have matched (**there**, **then**, **other**)
 - False positives
 - Not matching things that we should have matched (The)
 - False negatives

Exercise

- Task is to match “the”
 - 5 words: “the they theu The teo”
 - A solution: /the/
 - If the matches are “the they theu The teo”
- What are the false positives?
- What are the false negatives?



Errors

- Reducing the error rate for an application often involves two efforts:
 - Increasing accuracy, or precision, (minimizing **false positives**)
 - Increasing coverage, or recall, (minimizing **false negatives**).
 - Example: “They The the they”
 - /tʰe/ (They **The** the **they**)
 - / [tT]he/ (**They** The the **they**)
- We’ll be telling the same story for many tasks, in the whole semester

RE application examples

- XML tags
 - `Pattern.compile("<MeshHeadingList>.*</MeshHeadingList>", Pattern.DOTALL)`
 - `Pattern.compile("<DescriptorName[^<>]*>[^<>]*</DescriptorName>")`
- Time (e.g. 09:30): `^[0-2]?[0-3]:[0-5][0-9]$`

PPI extraction
step

A, TIMP-2 increases p27Kip1 association with Cdk4 and Cdk2.

`{interaction} := [interactor1] {interact-noun} {preposition} ([interactor2] | {protein-list});`

`{interact-noun} := association, co-localization, interaction ... ;`

`{preposition} := by, of, to, with ... ;`

`{protein-list} := ([interactor2]" , ")? [interactor2]" , "? and [interactor2];`


Finite state automata (FSA)

- Regular expressions
 - Compact textual strings (e.g. “/[tT]he/”)
 - Perfect for specifying patterns in programs or command-lines
 - Can be implemented as a FSA
- Finite state automata
 - Graphs
 - Can be described with a regular expression
 - A textual way of specifying the structure of FSA
 - FSA has a wide range of uses




FSA as Graphs

- Let's start with the sheep language from the text
 - /baa+!/
baa! baaa!
baaaa! ...



Sheep language FSA

- We can say the following things about this machine
 - It has 5 states
 - **b**, **a**, and **!** are in its alphabet
 - q_0 is the start state
 - q_4 is an accept state
 - It has 5 transitions



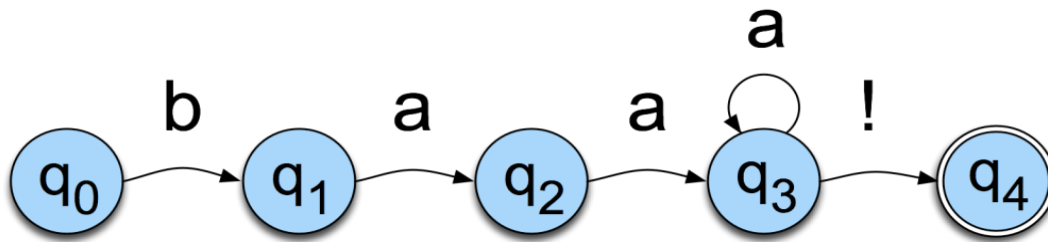
More Formally

- You can specify an FSA by enumerating the following things:
 - The set of states: Q
 - A finite alphabet: Σ
 - A start state
 - A set of accept/final states
 - A transition function that maps $Q \times \Sigma \rightarrow Q$
- Don't take term **alphabet** word too narrowly;
 - it just means we need **a finite set of symbols** in the input.



Yet Another View

- An FSA can ultimately be represented as tables



If you're in state 1 and you're looking at an input **a**, then go to state 2

	Input			
State	b	a	!	e
0	1			
1		2		
2		3		
3		3	4	
4:				



Recognition

- Recognition is the process of **determining if a string should be accepted** by a machine
- Or... it's the process of determining if a string is in the language we're defining with the machine
- Or... it's the process of **determining if a regular expression matches a string**




Recognition

- Traditionally, (Turing's notion) this process is depicted with a tape.





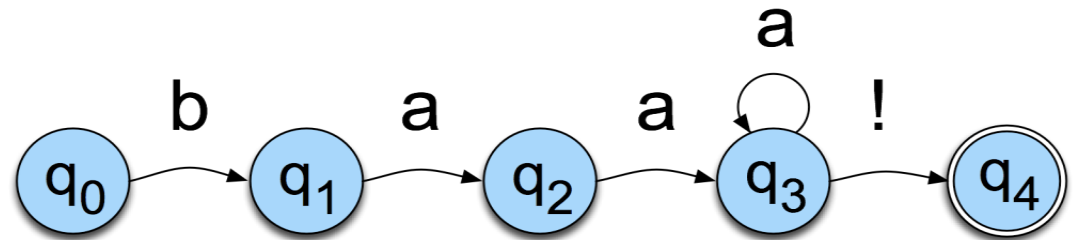
Recognition (D-Recognize)

- Simply a process of starting in the **start state**
- Examining the **current input**
- Consulting **the table**
- Going **to a new state** and updating the tape pointer.
- Until you run **out of tape**.
- Accept?



Example

	Input			
State	b	a	!	e
0	1			
1		2		
2		3		
3		3	4	
4:				



D-Recognize algorithm

function D-RECOGNIZE(*tape, machine*) **returns** accept or reject

index \leftarrow Beginning of tape

current-state \leftarrow Initial state of machine

loop

if End of input has been reached **then**

if *current-state* is an accept state **then**

return accept

else

return reject

elseif *transition-table*[*current-state*, *tape*[*index*]] is empty **then**

return reject

else

current-state \leftarrow *transition-table*[*current-state*, *tape*[*index*]]

index \leftarrow *index* + 1

end

Key Points

- Deterministic means that at each point in processing there is always one **unique** thing to do (no choices).
- D-recognize is a simple table-driven interpreter
- The algorithm is universal for all **unambiguous** regular languages.
 - To change the machine, you simply change the table.
- Matching strings with regular expressions (e.g., Perl, grep, etc.) is a matter of
 - translating the regular expression into a machine (a table) and
 - passing the table and the string to an interpreter that implement D-recognize

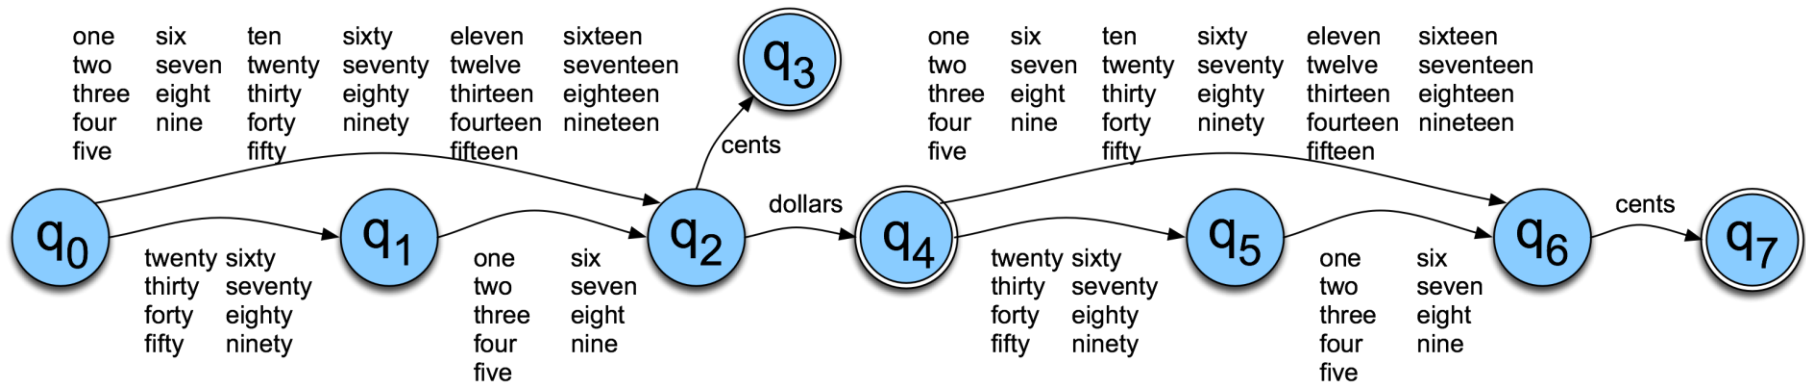


A short summary

- Regular expression
 - Very basic one /the/
 - More notations [], ?, *, .
- Two types of errors
 - false positives
 - false negatives
- Finite state automata
 - Representation
 - D-recognize algorithm to implement regular expression

Dollars and Cents

- Don't take term **alphabet** too narrowly; it just means we need a finite set of symbols in the input.




- Exercise: How to represent the expressions for phone numbers in Singapore with FSA/RE?

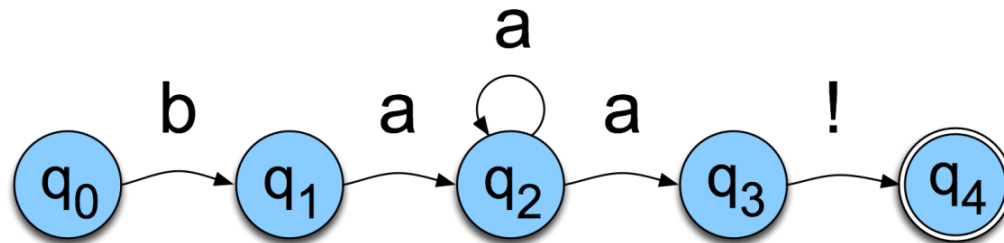
Non-Deterministic (NFSA)

- The sheep language from the text: **/baa+! /**

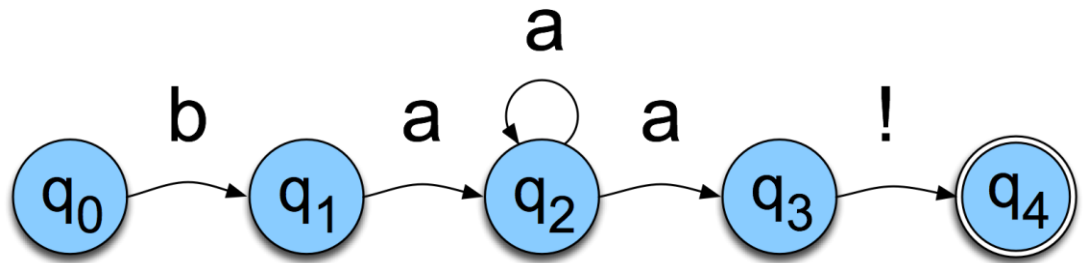
baa! baaa! baaaa! ...



- There are other machines that correspond to this same language



Yet Another View

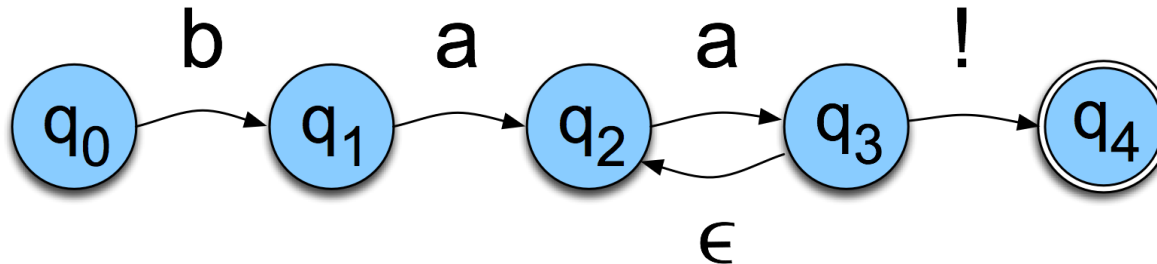


If you're in state 2 and you're looking at an input **a**, then go to state 2 or 3

	Input			
State	b	a	!	e
0	1			
1		2		
2		2, 3		
3			4	
4:				

Non-Deterministic (NFSA)

- Yet another technique
 - **Epsilon** transitions
 - These transitions do not examine or advance the tape during recognition



Equivalence

- Non-deterministic machines **can be converted to deterministic** ones with a fairly simple construction
- That means that they have the same power;
 - non-deterministic machines **are not more powerful than** deterministic ones in terms of the languages they can and can't characterize
- Non-determinism doesn't get us more formal power and it causes headaches so why bother?
 - More natural (understandable) solutions
 - Not always obvious to users whether or not the regex that they've produced is non-deterministic or not,
 - better to not make them worry about it

- In a ND FSA there exists at least one path through the machine for a string that is in the language defined by the machine.
- But **not all paths** directed through the machine for an accept string lead to an accept state.
- No paths through the machine lead to an accept state for a string not in the language.

Non-Deterministic Recognition

- Success in non-deterministic recognition occurs when a path is found through the machine that ends in an accept.
- Failure occurs when **all** of the possible paths for a given string lead to failure.

A summary

- Regular expression
 - Very basic regular expression /the/
 - More notations [], ?, *, .
- Two types of errors: false positive, false negative
- Finite state automata
 - Deterministic (FSA)
 - D-recognize algorithm to implement regular expression
 - Non-Deterministic (NFSA)
 - Two approaches to implementing regular expression

Readings

- Java.util.regex API
 - <http://docs.oracle.com/javase/8/docs/api/java/util/regex/package-summary.html>
- Java regexps tutorial
 - <http://docs.oracle.com/javase/tutorial/essential/regex/>
- RegExr: an online tool to learn, build, & test Regular Expressions
 - <http://regexr.com/>

